


Insurance Price Prediction with Linear Regression

In this notebook, I'm going to show how to make insurance price prediction with Linear Regression for new customers, using information such as their age, sex, BMI, children, smoking habits and region of residence.. After doing the project, will enter the patient's information such as age, sex, region into the model and then predict the patient's charge.



Data Description

The data set includes the following variables:

Age: Age of the Customer

Sex: Gender of the Customer.

bmi: A person's weight in kilograms divided by the square of height in meters.

children: The number of children of the customer.

smoker: If the insured person is a smoker or not.

region: Region where the customer lived.

charges: The premium of insurance.

Data Source: Open Source

Loading The Data

```
In [14]: import pandas as pd
```

```
In [15]: data = pd.read_csv("file:///C:/Users/angsh/OneDrive/Desktop/PRAXIS/Own%20Projects/ML/Regression%20project/insurance.c
```

```
In [16]: data.head(4)
```

Out[16]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061

Understanding The Data

```
In [17]: data.shape
```

```
Out[17]: (1338, 7)
```

```
In [18]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   age         1338 non-null   int64  
 1   sex         1338 non-null   object  
 2   bmi         1338 non-null   float64 
 3   children    1338 non-null   int64  
 4   smoker      1338 non-null   object  
 5   region      1338 non-null   object  
 6   charges     1338 non-null   float64 
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
```

Dealing with Missing Data

```
In [19]: data.isnull().sum()
```

```
Out[19]: age         0
sex           0
bmi           0
children      0
smoker        0
region        0
charges       0
dtype: int64
```

```
In [20]: data.dtypes
```

```
Out[20]: age           int64  
sex           object  
bmi           float64  
children      int64  
smoker        object  
region        object  
charges       float64  
dtype: object
```

Preprocessing The Data

Let's convert them to category type to use the model building later.

```
In [21]: data["sex"] = data["sex"].astype("category")  
data["smoker"] = data["smoker"].astype("category")  
data["region"] = data["region"].astype("category")
```

```
In [22]: data.dtypes
```

```
Out[22]: age           int64  
sex           category  
bmi           float64  
children      int64  
smoker        category  
region        category  
charges       float64  
dtype: object
```

Let's take a look at summary statistics.



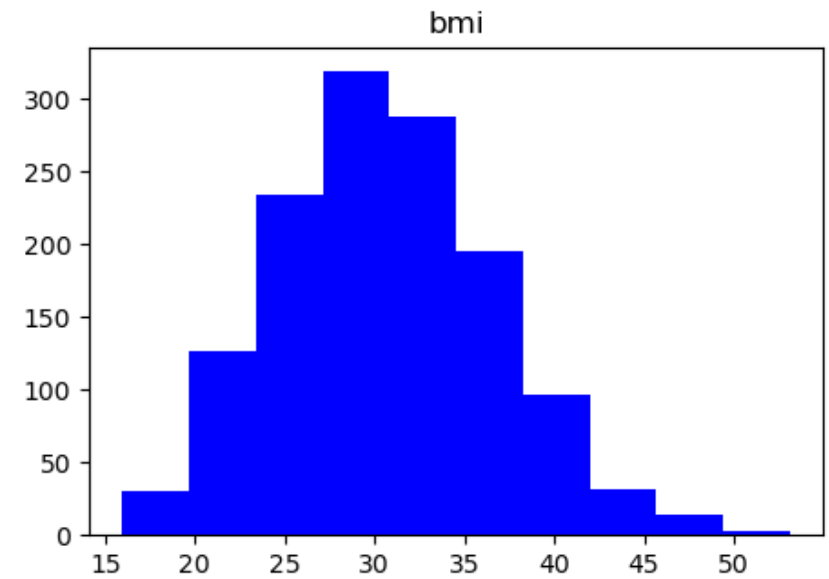

In [23]: `data.describe().T`

Out[23]:

	count	mean	std	min	25%	50%	75%	max
age	1338.0	39.207025	14.049960	18.0000	27.00000	39.000	51.000000	64.00000
bmi	1338.0	30.663397	6.098187	15.9600	26.29625	30.400	34.693750	53.13000
children	1338.0	1.094918	1.205493	0.0000	0.00000	1.000	2.000000	5.00000
charges	1338.0	13270.422265	12110.011237	1121.8739	4740.28715	9382.033	16639.912515	63770.42801

```
In [24]: data.hist(  
    ['age', 'bmi',  
     'children',  
     'charges'],figsize=(12,8),color='blue',grid=False)  
plt.show()
```

```
-----  
NameError                                Traceback (most recent call last)  
~\AppData\Local\Temp\ipykernel_13112\1812213538.py in <module>  
      3 'children',  
      4 'charges'],figsize=(12,8),color='blue',grid=False)  
----> 5 plt.show()  
  
NameError: name 'plt' is not defined
```




```
In [27]: smoke_data = data.groupby("smoker").mean().round(2)
```

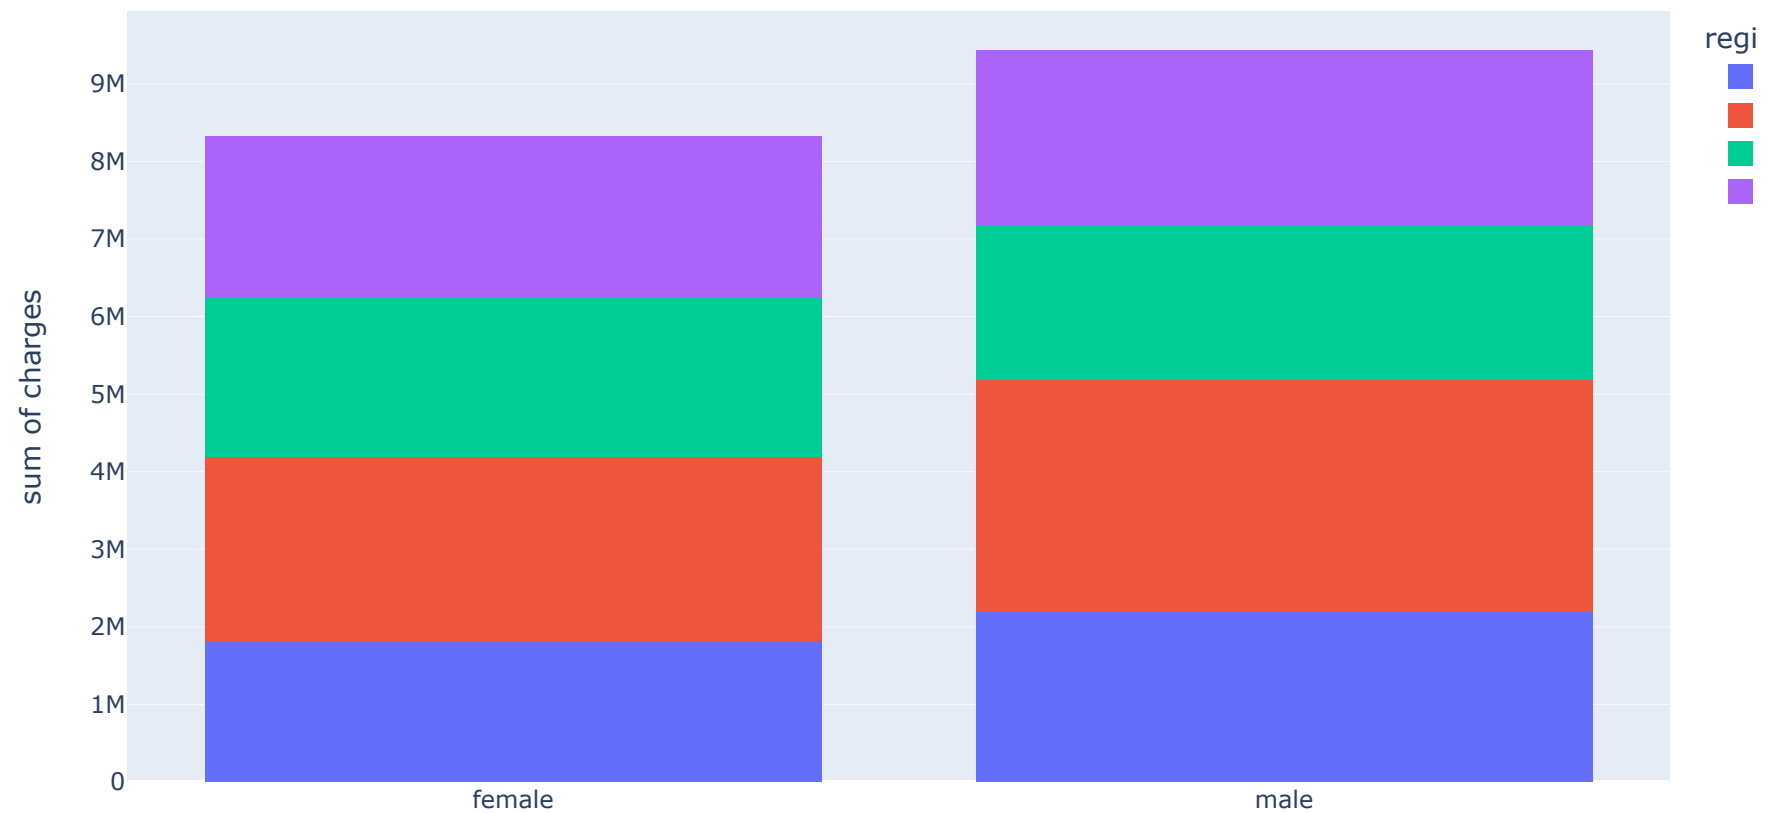
```
In [28]: smoke_data
```

Out[28]:

	age	bmi	children	charges
smoker				
no	39.39	30.65	1.09	8434.27
yes	38.51	30.71	1.11	32050.23

Preprocessing The Data


```
In [29]: import plotly.express as px  
px.histogram(data,x='sex',y = 'charges',color = 'region')
```



We can infer that from every region our customer base which has males are incurring more bills but interestingly females of northwest region are having more medical bills.

```
In [30]: fig = px.histogram(data,
                        x = 'charges',
                        marginal = "box",
                        color = 'children',
                        title = 'charges incurred by children')
fig.update_layout(bargap=0.1)
fig.show()
```

charges incurred by children



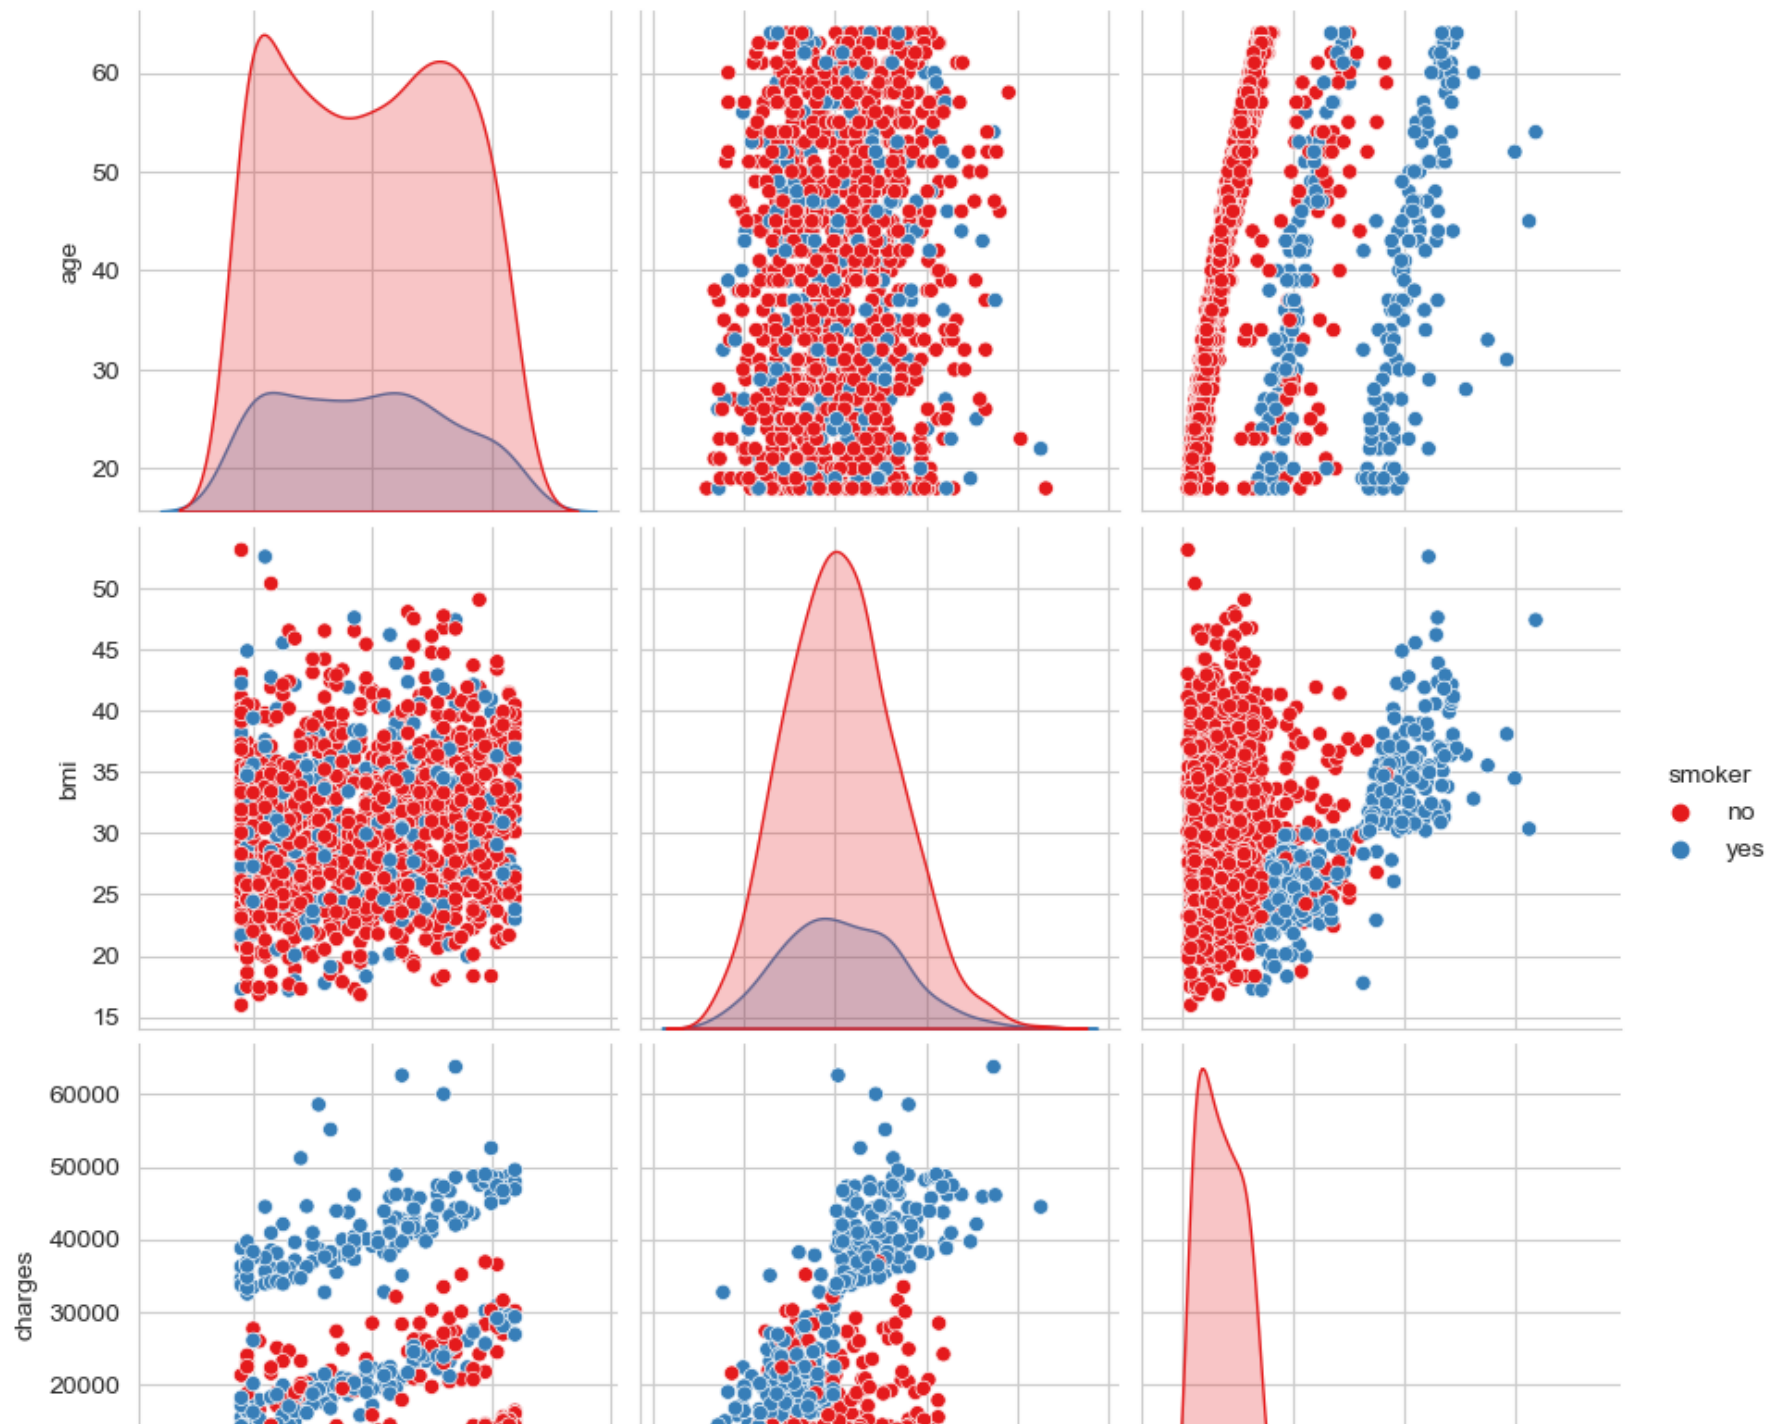
It seems that majority of our customers have 0 or 1 child and median charges vary between 8.5k to 11k dollars.


We can also conclude that people who have more children are given less priority in terms of pricing discounts.

```
In [31]: import seaborn as sns
```

```
In [32]: sns.set_style("whitegrid")
```


```
In [33]: import matplotlib.pyplot as plt
sns.pairplot(data[["age", "bmi", "charges", "smoker"]],
             hue = "smoker",
             height = 3,
             palette="Set1")
plt.show()
```



From the pairplot we can see that the charges to smokers is relatively high than non-smokers.

```
In [34]: sns.heatmap(data.corr(), annot = True)
plt.show()
```



One-Hot Encoding

Some categorical variables have subcategories such as sex and smoker. We need to convert these categorical variables into a form that the scikit-learn library can understand.

```
In [35]: data.columns
```

```
Out[35]: Index(['age', 'sex', 'bmi', 'children', 'smoker', 'region', 'charges'], dtype='object')
```

```
In [36]: data = pd.get_dummies(data)
```

```
In [37]: data.columns
```

```
Out[37]: Index(['age', 'bmi', 'children', 'charges', 'sex_female', 'sex_male',  
              'smoker_no', 'smoker_yes', 'region_northeast', 'region_northwest',  
              'region_southeast', 'region_southwest'],  
              dtype='object')
```

Creating Input and Output Variables

```
In [38]: X = data.drop("charges", axis = 1)
```

```
In [39]: y = data["charges"]
```

```
In [40]: target_range = y.max() - y.min()  
target_range
```

```
Out[40]: 62648.554110000005
```

Splitting The Dataset

```
In [41]: from sklearn.model_selection import train_test_split
```

```
In [61]: X_train, X_test, y_train, y_test = train_test_split(X, y, train_size= 0.80, random_state=1)
```

Building The Linear Regression Model

```
In [62]: from sklearn.linear_model import LinearRegression
```

```
In [63]: lr = LinearRegression()
```

```
In [64]: lr.fit(X_train, y_train)
```

```
Out[64]: LinearRegression()
```

Evaluating The Model

```
In [65]: lr.score(X_train, y_train).round(3)
```

```
Out[65]: 0.748
```

```
In [66]: lr.score(X_test, y_test).round(3)
```

```
Out[66]: 0.762
```

The model prediction is relatively near with its training set.

Now let's take a look at another metric, Root mean squared error, to evaluate the model.

```
In [67]: y_pred = lr.predict(X_test)
```

```
In [68]: from sklearn.metrics import mean_squared_error
```

```
In [69]: import math
```

```
In [70]: math.sqrt(mean_squared_error(y_test, y_pred))
```

```
Out[70]: 5956.454717976426
```

RMSE is relatively smaller in my linear regression model compared to the range of the target variable, which suggests that the model's predictive performance is reasonable.

```
In [71]: import numpy as np

# Train your model using X_train and y_train

# Make predictions on the test data
y_pred = lr.predict(X_test)

# Calculate RMSLE
log_y_pred = np.log1p(y_pred) # Take the Logarithm of the predicted values
log_y_test = np.log1p(y_test) # Take the Logarithm of the actual values
squared_diff = (log_y_pred - log_y_test) ** 2 # Calculate squared difference
mean_squared_log_error = np.mean(squared_diff) # Average of squared differences
rmsle = np.sqrt(mean_squared_log_error) # Square root of the average

print("RMSLE:", rmsle)
```

RMSLE: 0.49544319054018776

C:\Users\angsh\AppData\Local\Temp\ipykernel_13112\1292476574.py:9: RuntimeWarning:

invalid value encountered in log1p

Now try to make my model prediction more good with Random Forest Regressor.

Building The Random Forest Model

```
In [72]: import numpy as np
from sklearn.ensemble import RandomForestRegressor
# Create and train the Random Forest model
rf = RandomForestRegressor(n_estimators=100, random_state=42)
rf.fit(X_train, y_train)

# Predict on the test set
y_pred = rf.predict(X_test)

# Calculate the root mean squared error (RMSE)
mse=mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
print("RMSE:", rmse)
```

RMSE: 4655.54818633615

Here RMSE is relatively more smaller than the linear regression model , which suggests that the model's predictive performance improved with it than LR.

```
In [73]: import numpy as np

# Train your model using X_train and y_train

# Make predictions on the test data
y_pred = rf.predict(X_test)

# Calculate RMSLE
log_y_pred = np.log1p(y_pred) # Take the Logarithm of the predicted values
log_y_test = np.log1p(y_test) # Take the Logarithm of the actual values
squared_diff = (log_y_pred - log_y_test) ** 2 # Calculate squared difference
mean_squared_log_error = np.mean(squared_diff) # Average of squared differences
rmsle = np.sqrt(mean_squared_log_error) # Square root of the average

print("RMSLE:", rmsle)
```

RMSLE: 0.43228881831505095

Evaluating The Model

```
In [98]: rf.score(X_train, y_train).round(3)
```

```
Out[98]: 0.976
```

```
In [99]: rf.score(X_test, y_test).round(3)
```

```
Out[99]: 0.855
```

The Random Forest model score got better than Linear Regression

Building The KNR Model

```
In [100]: from sklearn.neighbors import KNeighborsRegressor
from sklearn.metrics import mean_squared_error

# Step 1: Import the necessary Libraries

# Step 2: Create an instance of the KNeighborsRegressor
knr = KNeighborsRegressor(n_neighbors=5)

# Step 3: Fit the KNR model to your training data
knr.fit(X_train, y_train)

# Step 4: Use the trained model to make predictions on the test data
y_pred = knr.predict(X_test)

# Step 5: Evaluate the performance of the KNR model using Mean Squared Error (MSE)
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
print("RMSE:", rmse)
```

RMSE: 10335.483390128351

```
In [110]: import numpy as np

# Train your model using X_train and y_train

# Make predictions on the test data
y_pred = knr.predict(X_test)

# Calculate RMSLE
log_y_pred = np.log1p(y_pred) # Take the Logarithm of the predicted values
log_y_test = np.log1p(y_test) # Take the Logarithm of the actual values
squared_diff = (log_y_pred - log_y_test) ** 2 # Calculate squared difference
mean_squared_log_error = np.mean(squared_diff) # Average of squared differences
rmsle = np.sqrt(mean_squared_log_error) # Square root of the average

print("RMSLE:", rmsle)
```

RMSLE: 0.7451786340817319

Predicting a New Data

I'm going to predict the first row as an example.

```
In [101]: data_new = X_train[:4]
```

```
In [102]: data_new
```

```
Out[102]:
```

	age	bmi	children	sex_female	sex_male	smoker_no	smoker_yes	region_northeast	region_northwest	region_southeast	region_southwest
216	53	26.600	0	1	0	1	0	0	1	0	
731	53	21.400	1	0	1	1	0	0	0	0	
866	18	37.290	0	0	1	1	0	0	0	1	
202	60	24.035	0	1	0	1	0	0	1	0	

Predict with Linear Regressor

```
In [103]: lr.predict(data_new)
```

```
Out[103]: array([10508.41885042, 8494.95651816, 4049.96586839, 11485.89042227])
```

Predict with Random Forest Regressor

```
In [104]: rf.predict(data_new)
```

```
Out[104]: array([11780.1045341, 11407.0849706, 1325.3675585, 14226.998494 ])
```

Predict with K- Nearest Regressor

```
In [105]: knr.predict(data_new)
```

```
Out[105]: array([13657.71103 , 10217.50862 , 10937.09363 , 18416.929286])
```

So let's see the real value.

```
In [106]: y_train[:4]
```

```
Out[106]: 216    10355.64100  
731    10065.41300  
866     1141.44510  
202    13012.20865  
Name: charges, dtype: float64
```

As you can see, that out of all the models , Random Forest is giving the most accurate result compare to others. And Hence we can say that Random Forest is the best regressor to us to work on this dataset.

```
In [ ]:
```